

movie_recommendation

December 8, 2024

```
[5]: %matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from ast import literal_eval
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.metrics.pairwise import linear_kernel, cosine_similarity
from nltk.stem.snowball import SnowballStemmer
from nltk.stem.wordnet import WordNetLemmatizer
from nltk.corpus import wordnet
from surprise import Reader, Dataset, SVD
from surprise.model_selection import cross_validate

import warnings; warnings.simplefilter('ignore')
```

```
[6]: f_path=r'/Users/91950/Documents/293C-Add/movies_metadata.csv'
md = pd.read_csv(f_path)
md.head()
```

```
[6]:      adult      belongs_to_collection      budget \
0  False  {'id': 10194, 'name': 'Toy Story Collection', ...  30000000
1  False                                     NaN  65000000
2  False  {'id': 119050, 'name': 'Grumpy Old Men Collect...      0
3  False                                     NaN  16000000
4  False  {'id': 96871, 'name': 'Father of the Bride Col...      0

      genres \
0  [{'id': 16, 'name': 'Animation'}, {'id': 35, 'name': 'Comedy'}]
1  [{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}]
2  [{'id': 10749, 'name': 'Romance'}, {'id': 35, 'name': 'Comedy'}]
3  [{'id': 35, 'name': 'Comedy'}, {'id': 18, 'name': 'Drama'}]
4  [{'id': 35, 'name': 'Comedy'}]

      homepage      id      imdb_id  original_language \
0  http://toystory.disney.com/toy-story      862  tt0114709      en
1      NaN      8844  tt0113497      en
```

2		NaN	15602	tt0113228	en
3		NaN	31357	tt0114885	en
4		NaN	11862	tt0113041	en

	original_title \
0	Toy Story
1	Jumanji
2	Grumpier Old Men
3	Waiting to Exhale
4	Father of the Bride Part II

	overview ...	release_date \
0	Led by Woody, Andy's toys live happily in his ...	1995-10-30
1	When siblings Judy and Peter discover an encha...	1995-12-15
2	A family wedding reignites the ancient feud be...	1995-12-22
3	Cheated on, mistreated and stepped on, the wom...	1995-12-22
4	Just when George Banks has recovered from his ...	1995-02-10

	revenue	runtime	spoken_languages \
0	373554033.0	81.0	[{'iso_639_1': 'en', 'name': 'English'}]
1	262797249.0	104.0	[{'iso_639_1': 'en', 'name': 'English'}, {'iso...
2	0.0	101.0	[{'iso_639_1': 'en', 'name': 'English'}]
3	81452156.0	127.0	[{'iso_639_1': 'en', 'name': 'English'}]
4	76578911.0	106.0	[{'iso_639_1': 'en', 'name': 'English'}]

	status	tagline \
0	Released	NaN
1	Released	Roll the dice and unleash the excitement!
2	Released	Still Yelling. Still Fighting. Still Ready for...
3	Released	Friends are the people who let you be yourself...
4	Released	Just When His World Is Back To Normal... He's ...

	title	video	vote_average	vote_count
0	Toy Story	False	7.7	5415.0
1	Jumanji	False	6.9	2413.0
2	Grumpier Old Men	False	6.5	92.0
3	Waiting to Exhale	False	6.1	34.0
4	Father of the Bride Part II	False	5.7	173.0

[5 rows x 24 columns]

```
[7]: md['genres'] = md['genres'].fillna('').apply(literal_eval).apply(lambda x: [
    ↪ [i['name'] for i in x] if isinstance(x, list) else [])
```

```
[8]: vote_counts = md[md['vote_count'].notnull()]['vote_count'].astype('int')
vote_averages = md[md['vote_average'].notnull()]['vote_average'].astype('int')
C = vote_averages.mean()
```

```
C
```

```
[8]: 5.244896612406511
```

```
[9]: m = vote_counts.quantile(0.95)
m
```

```
[9]: 434.0
```

```
[10]: md['year'] = pd.to_datetime(md['release_date'], errors='coerce').apply(lambda x:
    ↪ str(x).split('-')[0] if x != np.nan else np.nan)
```

```
[11]: qualified = md[(md['vote_count'] >= m) & (md['vote_count'].notnull()) &
    ↪ (md['vote_average'].notnull())][['title', 'year', 'vote_count',
    ↪ 'vote_average', 'popularity', 'genres']]
qualified['vote_count'] = qualified['vote_count'].astype('int')
qualified['vote_average'] = qualified['vote_average'].astype('int')
qualified.shape
```

```
[11]: (2274, 6)
```

```
[12]: def weighted_rating(x):
    v = x['vote_count']
    R = x['vote_average']
    return (v/(v+m) * R) + (m/(m+v) * C)
```

```
[13]: qualified['wr'] = qualified.apply(weighted_rating, axis=1)
```

```
[14]: qualified = qualified.sort_values('wr', ascending=False).head(250)
```

```
[15]: qualified.head(15)
```

```
[15]:
```

	title	year	vote_count	\
15480	Inception	2010	14075	
12481	The Dark Knight	2008	12269	
22879	Interstellar	2014	11187	
2843	Fight Club	1999	9678	
4863	The Lord of the Rings: The Fellowship of the Ring	2001	8892	
292	Pulp Fiction	1994	8670	
314	The Shawshank Redemption	1994	8358	
7000	The Lord of the Rings: The Return of the King	2003	8226	
351	Forrest Gump	1994	8147	
5814	The Lord of the Rings: The Two Towers	2002	7641	
256	Star Wars	1977	6778	
1225	Back to the Future	1985	6239	
834	The Godfather	1972	6024	
1154	The Empire Strikes Back	1980	5998	
46	Se7en	1995	5915	

	vote_average	popularity \
15480	8	29.108149
12481	8	123.167259
22879	8	32.213481
2843	8	63.869599
4863	8	32.070725
292	8	140.950236
314	8	51.645403
7000	8	29.324358
351	8	48.307194
5814	8	29.423537
256	8	42.149697
1225	8	25.778509
834	8	41.109264
1154	8	19.470959
46	8	18.45743

	genres	wr
15480	[Action, Thriller, Science Fiction, Mystery, A...	7.917588
12481	[Drama, Action, Crime, Thriller]	7.905871
22879	[Adventure, Drama, Science Fiction]	7.897107
2843	[Drama]	7.881753
4863	[Adventure, Fantasy, Action]	7.871787
292	[Thriller, Crime]	7.868660
314	[Drama, Crime]	7.864000
7000	[Adventure, Fantasy, Action]	7.861927
351	[Comedy, Drama, Romance]	7.860656
5814	[Adventure, Fantasy, Action]	7.851924
256	[Adventure, Action, Science Fiction]	7.834205
1225	[Adventure, Comedy, Science Fiction, Family]	7.820813
834	[Drama, Crime]	7.814847
1154	[Adventure, Action, Science Fiction]	7.814099
46	[Crime, Mystery, Thriller]	7.811669

```
[16]: s = md.apply(lambda x: pd.Series(x['genres']),axis=1).stack().
      ↪reset_index(level=1, drop=True)
      s.name = 'genre'
      gen_md = md.drop('genres', axis=1).join(s)
```

```
[19]: import matplotlib.pyplot as plt
      import pandas as pd

      def build_chart(genre, percentile=0.85):
          # Filter the dataframe for the selected genre
          df = gen_md[gen_md['genre'] == genre]
```

```

# Extract necessary columns and clean data
vote_counts = df[df['vote_count'].notnull()]['vote_count'].astype('int')
vote_averages = df[df['vote_average'].notnull()]['vote_average'].
↳astype('int')
C = vote_averages.mean()
m = vote_counts.quantile(percentile)

# Select qualified movies based on vote count and vote average thresholds
qualified = df[(df['vote_count'] >= m) & (df['vote_count'].notnull()) &
↳(df['vote_average'].notnull())][['title', 'year', 'vote_count',
↳'vote_average', 'popularity']]
qualified['vote_count'] = qualified['vote_count'].astype('int')
qualified['vote_average'] = qualified['vote_average'].astype('int')

# Calculate weighted rating (wr)
qualified['wr'] = qualified.apply(lambda x: (x['vote_count']/
↳(x['vote_count']+m) * x['vote_average']) + (m/(m+x['vote_count'])) * C),
↳axis=1)

# Sort the movies by weighted rating and select top 250
qualified = qualified.sort_values('wr', ascending=False).head(250)

# Ensure the 'popularity' column is numeric, and drop rows with NaN values
↳in popularity
qualified['popularity'] = pd.to_numeric(qualified['popularity'],
↳errors='coerce')
qualified = qualified.dropna(subset=['popularity'])

# Create a pie chart based on 'popularity' (you can choose another feature
↳like 'vote_average' or 'wr' if desired)
popularity_bins = ['Low', 'Medium', 'High']
# Define thresholds for categorizing popularity
qualified['popularity_category'] = pd.cut(qualified['popularity'], bins=[0,
↳20, 60, 100], labels=popularity_bins)

# Get the count of each popularity category
popularity_counts = qualified['popularity_category'].value_counts()


# Plot the pie chart with multiple colors
plt.figure(figsize=(7, 7))
plt.pie(popularity_counts, labels=popularity_counts.index, autopct='%1.
↳1f%%', colors=plt.cm.Paired.colors)
plt.title(f'Popularity Distribution of {genre} Movies (Top 250 by Weighted
↳Rating)', fontsize=14)
plt.show()

```

```
# Return the filtered DataFrame
return qualified
```

```
[20]: build_chart('Drama').head(25)
```

Popularity Distribution of Drama Movies (Top 250 by Weighted Rating)



```
[20]:
```

	title	year	vote_count	vote_average	\
10309	Dilwale Dulhania Le Jayenge	1995	661	9	
12481	The Dark Knight	2008	12269	8	
22879	Interstellar	2014	11187	8	
2843	Fight Club	1999	9678	8	
314	The Shawshank Redemption	1994	8358	8	
351	Forrest Gump	1994	8147	8	
834	The Godfather	1972	6024	8	
24860	The Imitation Game	2014	5895	8	
359	The Lion King	1994	5520	8	
18465	The Intouchables	2011	5410	8	
22841	The Grand Budapest Hotel	2014	4644	8	
586	The Silence of the Lambs	1991	4549	8	

11354	The Prestige	2006	4510	8
522	Schindler's List	1993	4436	8
23673	Whiplash	2014	4376	8
289	Leon: The Professional	1994	4293	8
3030	The Green Mile	1999	4166	8
2211	Life Is Beautiful	1997	3643	8
1163	A Clockwork Orange	1971	3432	8
1178	The Godfather: Part II	1974	3418	8
49	The Usual Suspects	1995	3334	8
1170	GoodFellas	1990	3211	8
2216	American History X	1998	3120	8
4135	Scarface	1983	3017	8
1152	One Flew Over the Cuckoo's Nest	1975	3001	8

	popularity	wr	popularity_category
10309	34.457024	8.607205	Medium
12481	123.167259	7.983111	NaN
22879	32.213481	7.981489	Medium
2843	63.869599	7.978628	High
314	51.645403	7.975286	Medium
351	48.307194	7.974653	Medium
834	41.109264	7.965843	Medium
24860	31.595940	7.965106	Medium
359	21.605761	7.962771	Medium
18465	16.086919	7.962025	Low
22841	14.442048	7.955873	Low
586	4.307222	7.954968	Low
11354	16.945560	7.954586	Low
522	41.725123	7.953842	Medium
23673	64.299990	7.953222	High
289	20.477329	7.952334	Medium
3030	19.966780	7.950910	Low
2211	39.394970	7.944021	Medium
1163	17.112594	7.940662	Low
1178	36.629307	7.940425	Medium
49	16.302466	7.938961	Low
1170	15.424092	7.936682	Low
2216	18.157166	7.934884	Low
4135	11.299673	7.932721	Low
1152	35.529554	7.932372	Medium

```
[25]: f_path=r'/Users/91950/Documents/293C-Add/links_small.csv'
links_small = pd.read_csv(f_path)
links_small = links_small[links_small['tmdbId'].notnull()]['tmdbId'].
↳ astype('int')
```

```
[26]: md = md.drop([19730, 29503, 35587])
```

```
[29]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Assuming md is a DataFrame and links_small is a list or array of IDs
md['id'] = md['id'].astype('int')

# Filter the dataset based on 'id' in 'links_small'
smd = md[md['id'].isin(links_small)]

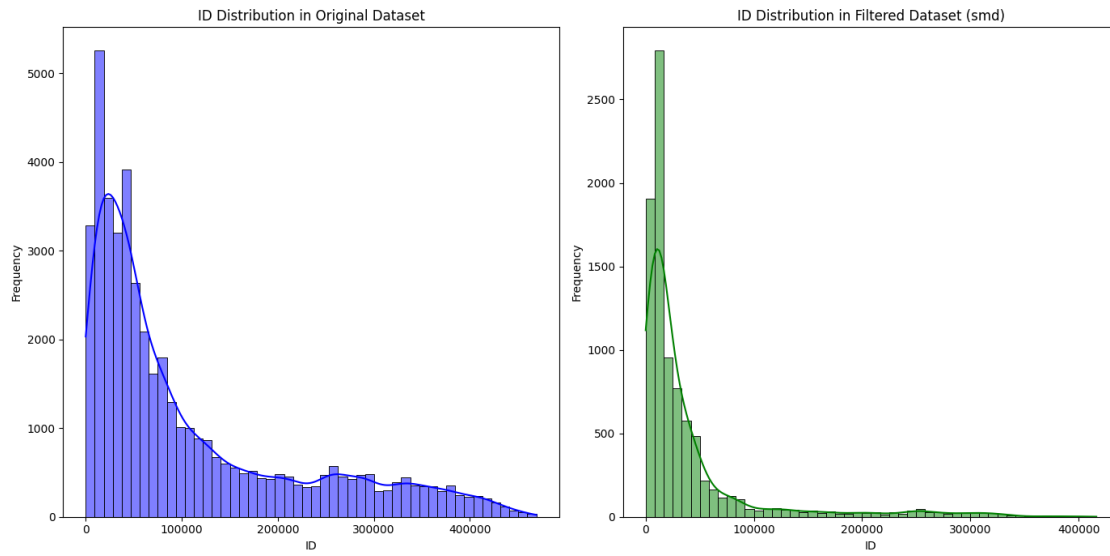
# Plot the distribution of 'id' in the original dataset and the filtered one
plt.figure(figsize=(14, 7))

# Plot the distribution of 'id' in the original dataset
plt.subplot(1, 2, 1)
sns.histplot(md['id'], bins=50, kde=True, color='blue')
plt.title("ID Distribution in Original Dataset")
plt.xlabel("ID")
plt.ylabel("Frequency")

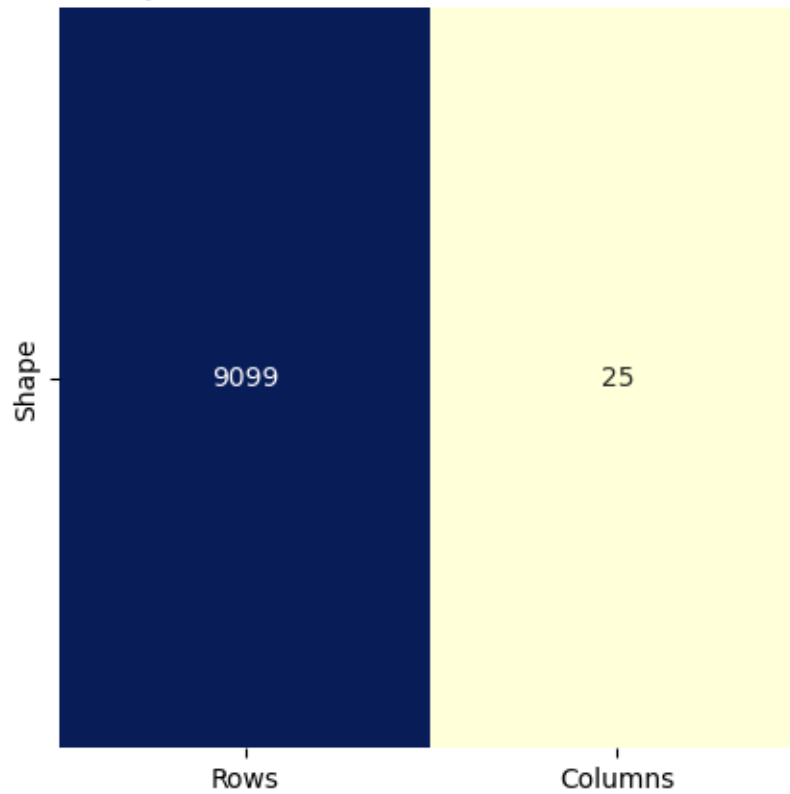
# Plot the distribution of 'id' in the filtered dataset (smd)
plt.subplot(1, 2, 2)
sns.histplot(smd['id'], bins=50, kde=True, color='green')
plt.title("ID Distribution in Filtered Dataset (smd)")
plt.xlabel("ID")
plt.ylabel("Frequency")

plt.tight_layout()
plt.show()

# Now print the shape of the filtered dataset as a graph representation
smd_shape = smd.shape
plt.figure(figsize=(5, 5))
sns.heatmap([[smd_shape[0], smd_shape[1]]], annot=True, fmt='d', cmap="YlGnBu",
            cbar=False,
            xticklabels=["Rows", "Columns"], yticklabels=["Shape"])
plt.title(f"Shape of Filtered Dataset (smd): {smd_shape}")
plt.show()
```

Shape of Filtered Dataset (smd): (9099, 25)



```
[37]: smd['tagline'] = smd['tagline'].fillna('')
smd['description'] = smd['overview'] + smd['tagline']
smd['description'] = smd['description'].fillna('')
tf = TfidfVectorizer(analyzer='word', ngram_range=(1, 2), min_df=0.01,
    ↪stop_words='english')
tfidf_matrix = tf.fit_transform(smd['description'])
cosine_sim = linear_kernel(tfidf_matrix, tfidf_matrix)
cosine_sim[0]
```

```
[37]: array([1.          , 0.          , 0.          , ..., 0.          , 0.06566461,
        0.          ])
```

```
[38]: smd = smd.reset_index()
titles = smd['title']
indices = pd.Series(smd.index, index=smd['title'])
```

```
[53]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

def get_recommendations_with_graph(title, top_n=10):
    # Get the index of the movie that matches the title
    idx = indices[title]

    # Get the similarity scores for the movie
    sim_scores = list(enumerate(cosine_sim[idx]))

    # Sort the movies based on similarity scores
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

    # Get the top N similar movie indices (excluding the original movie)
    sim_scores = sim_scores[1:top_n+1]
    movie_indices = [i[0] for i in sim_scores]

    # Get the corresponding movie titles and similarity scores
    recommended_titles = titles.iloc[movie_indices]
    similarity_scores = [x[1] for x in sim_scores]

    # Create a DataFrame for easier manipulation
    recommendations_df = pd.DataFrame({
        'Movie Title': recommended_titles,
        'Similarity Score': similarity_scores
    })

    # Set up the plot
    plt.figure(figsize=(10, 6))
```

```

# Create a horizontal bar plot using seaborn
sns.barplot(x='Similarity Score', y='Movie Title', data=recommendations_df,
palette='viridis')

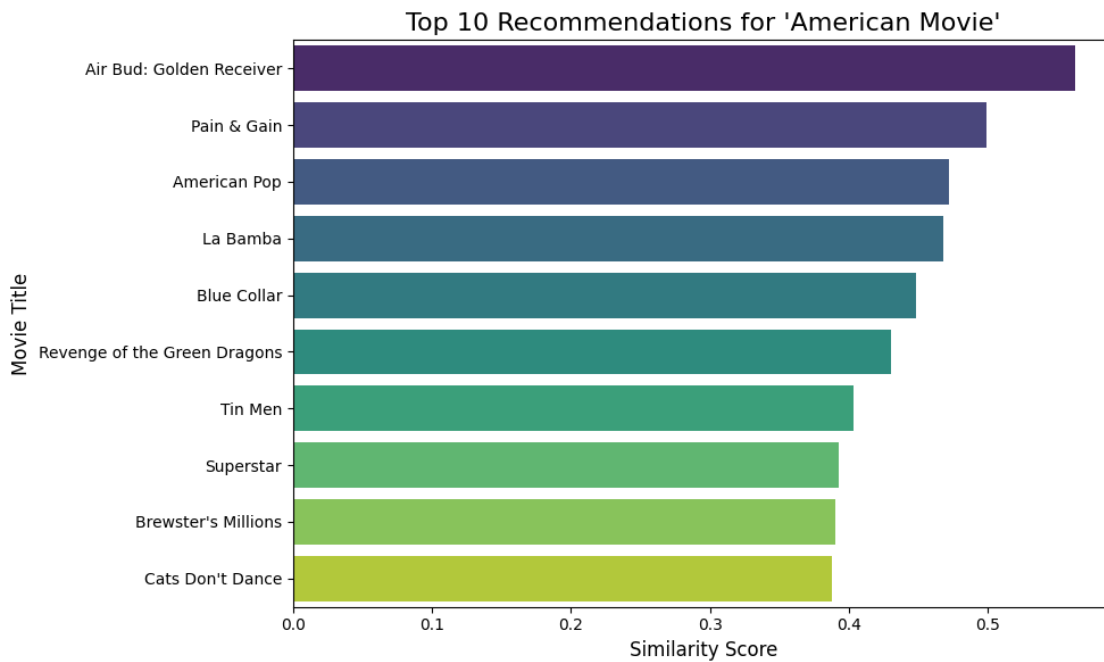
# Title and labels
plt.title(f"Top {top_n} Recommendations for '{title}'", fontsize=16)
plt.xlabel('Similarity Score', fontsize=12)
plt.ylabel('Movie Title', fontsize=12)

# Show the plot
plt.tight_layout()
plt.show()

return recommendations_df

# Example usage
get_recommendations_with_graph('American Movie')

```



```

[53]:
1707      Air Bud: Golden Receiver      0.562491
8283              Pain & Gain      0.499357
2973      American Pop      0.472206
2773              La Bamba      0.467615
2645      Blue Collar      0.448342
8623  Revenge of the Green Dragons      0.430133

```

1129	Tin Men	0.403227
2329	Superstar	0.392529
3208	Brewster's Millions	0.389989
1198	Cats Don't Dance	0.388185

```
[64]: from sklearn.model_selection import train_test_split
import pandas as pd

# Load the data
f_path = r'/Users/91950/Documents/293C-Add/ratings_small.csv'
ratings = pd.read_csv(f_path)

# Split the data into train and test sets using sklearn's train_test_split
train_df, test_df = train_test_split(ratings, test_size=0.2, random_state=42)

# Setup Reader and load the dataset
reader = Reader()
train_data = Dataset.load_from_df(train_df[['userId', 'movieId', 'rating']],
    ↪reader)
test_data = Dataset.load_from_df(test_df[['userId', 'movieId', 'rating']],
    ↪reader)

# Check the first few rows of the ratings data
ratings.head()
```

```
[64]:  userId  movieId  rating  timestamp
0      1      31      2.5  1260759144
1      1     1029      3.0  1260759179
2      1     1061      3.0  1260759182
3      1     1129      2.0  1260759185
4      1     1172      4.0  1260759205
```

```
[66]: svd = SVD()
results = cross_validate(svd, data, measures=['RMSE', 'MAE'], cv=5,
    ↪verbose=True)
```

Evaluating RMSE, MAE of algorithm SVD on 5 split(s).

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	0.8939	0.8967	0.8951	0.9005	0.8903	0.8953	0.0033
MAE (testset)	0.6869	0.6907	0.6895	0.6930	0.6843	0.6889	0.0030
Fit time	1.28	1.63	1.62	1.34	1.34	1.44	0.15
Test time	0.48	0.25	0.20	0.16	0.13	0.24	0.12

```
[68]: trainset = data.build_full_trainset()
svd.fit(trainset)
```

```
[68]: <surprise.prediction_algorithms.matrix_factorization.SVD at 0x1aa91f8c830>
```

```
[69]: ratings[ratings['userId'] == 1]
```

```
[69]:
```

	userId	movieId	rating	timestamp
0	1	31	2.5	1260759144
1	1	1029	3.0	1260759179
2	1	1061	3.0	1260759182
3	1	1129	2.0	1260759185
4	1	1172	4.0	1260759205
5	1	1263	2.0	1260759151
6	1	1287	2.0	1260759187
7	1	1293	2.0	1260759148
8	1	1339	3.5	1260759125
9	1	1343	2.0	1260759131
10	1	1371	2.5	1260759135
11	1	1405	1.0	1260759203
12	1	1953	4.0	1260759191
13	1	2105	4.0	1260759139
14	1	2150	3.0	1260759194
15	1	2193	2.0	1260759198
16	1	2294	2.0	1260759108
17	1	2455	2.5	1260759113
18	1	2968	1.0	1260759200
19	1	3671	3.0	1260759117

```
[70]: svd.predict(1, 302, 3)
```

```
[70]: Prediction(uid=1, iid=302, r_ui=3, est=2.8230487751680573,  
details={'was_impossible': False})
```

```
[71]: import matplotlib.pyplot as plt  
import seaborn as sns  
from surprise import SVD, Dataset, Reader  
from surprise.model_selection import cross_validate  
import pandas as pd  
  
# Load the data  
f_path = r'/Users/91950/Documents/293C-Add/ratings_small.csv'  
ratings = pd.read_csv(f_path)  
  
# Setup Reader and load the dataset  
reader = Reader()  
data = Dataset.load_from_df(ratings[['userId', 'movieId', 'rating']], reader)  
  
# Create an SVD model  
svd = SVD()  
  
# Evaluate the model using cross-validation
```

```

results = cross_validate(svd, data, measures=['RMSE', 'MAE'], cv=5,
    verbose=True)

# Plot RMSE and MAE results from cross-validation
cv_results = pd.DataFrame(results)
plt.figure(figsize=(12, 6))

# Plot RMSE
plt.subplot(1, 2, 1)
sns.boxplot(data=cv_results['test_rmse'])
plt.title('Cross-validation RMSE')
plt.ylabel('RMSE')

# Plot MAE
plt.subplot(1, 2, 2)
sns.boxplot(data=cv_results['test_mae'])
plt.title('Cross-validation MAE')
plt.ylabel('MAE')

plt.tight_layout()
plt.show()

# Train the model on the full training set
trainset = data.build_full_trainset()
svd.fit(trainset)


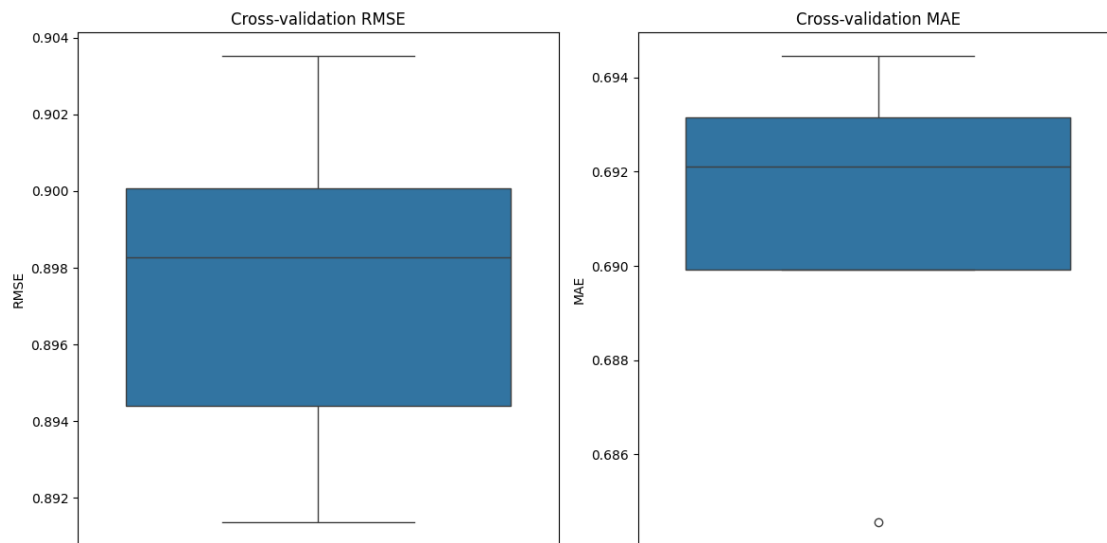
# Show ratings for userId 1
user_ratings = ratings[ratings['userId'] == 1]

# Plot a bar chart of ratings for userId 1
plt.figure(figsize=(10, 6))
sns.barplot(x='movieId', y='rating', data=user_ratings, palette='viridis')
plt.title('Ratings given by User 1')
plt.xlabel('Movie ID')
plt.ylabel('Rating')
plt.xticks(rotation=90)
plt.show()

```

Evaluating RMSE, MAE of algorithm SVD on 5 split(s).

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	0.8914	0.9035	0.8983	0.8944	0.9001	0.8975	0.0043
MAE (testset)	0.6846	0.6945	0.6921	0.6899	0.6932	0.6908	0.0035
Fit time	1.31	1.26	1.14	1.28	1.24	1.25	0.06
Test time	0.17	0.13	0.16	0.16	0.42	0.21	0.11



[]: